

Transfer_Learning

October 21, 2019

1 Trabalho 1

1.1 OCR - Optical Character Recognizer

Aluno: Bruno Flávio Ferreira - 9791330

Na primeira parte do trabalho, foi utilizada a técnica de Transfer Learning, para treinar uma rede neural capaz de reconhecer caracteres.

```
[0]: from __future__ import absolute_import, division, print_function, \
      ↪ unicode_literals
import matplotlib.pyplot as plt
import tensorflow as tf
import scipy.io as io
import tensorflow_hub as hub
from shutil import copyfile
from tensorflow.keras import layers
import pandas as pd
import PIL.Image as Image
import os
import numpy as np
import time
```

Primeiro, carregamos o dataset Chars74K, 75% é carregado como dataset de treinamento, 25% como validação

```
[7]: IMAGE_SHAPE = (224, 224)
dataset_folder = './TrainSet'
image_generator = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1/
↪ 255, validation_split=0.25)
image_data = image_generator.flow_from_directory(dataset_folder, \
↪ target_size=IMAGE_SHAPE, subset='training')
image_validation = image_generator.flow_from_directory(dataset_folder, \
↪ target_size=IMAGE_SHAPE, subset='validation')
```

Found 5806 images belonging to 62 classes.

Found 1899 images belonging to 62 classes.

A Rede que será utilizada para o Transfer Learning é a Inception V3 da Google.

Primeiramente, foi carregado o Inception V3 e executado sobre algumas imagens do dataset para vermos as predições antes de realizarmos a técnica do Transfer Learning

```
[8]: classifier_url = "https://tfhub.dev/google/tf2-preview/inception_v3/
    ↳ classification/4" #@param {type: "string"}
labels_path = tf.keras.utils.get_file('ImageNetLabels.txt', 'https://storage.
    ↳ googleapis.com/download.tensorflow.org/data/ImageNetLabels.txt')
imagenet_labels = np.array(open(labels_path).read().splitlines())
```

Downloading data from

https://storage.googleapis.com/download.tensorflow.org/data/ImageNetLabels.txt
16384/10484 [=====] - 0s 0us/step

```
[9]: classifier = tf.keras.Sequential([
    hub.KerasLayer(classifier_url, input_shape=IMAGE_SHAPE+(3,))
])

for image_batch, label_batch in image_data:
    print("Image batch shape: ", image_batch.shape)
    print("Label batch shape: ", label_batch.shape)
    break

result_batch = classifier.predict(image_batch)
print(result_batch.shape)

predicted_class_names = imagenet_labels[np.argmax(result_batch, axis=-1)]

plt.figure(figsize=(10,9))
plt.subplots_adjust(hspace=0.5)
for n in range(30):
    plt.subplot(6,5,n+1)
    plt.imshow(image_batch[n])
    plt.title(predicted_class_names[n])
    plt.axis('off')
    _ = plt.suptitle("ImageNet predictions")

plt.show()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/ops/resource_variable_ops.py:1781: calling BaseResourceVariable.__init__ (from tensorflow.python.ops.resource_variable_ops) with constraint is deprecated and will be removed in a future version.

Instructions for updating:

If using Keras pass *_constraint arguments to layers.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/ops/resource_variable_ops.py:1781: calling BaseResourceVariable.__init__ (from tensorflow.python.ops.resource_variable_ops)

with constraint is deprecated and will be removed in a future version.

Instructions for updating:


If using Keras pass *_constraint arguments to layers.

Image batch shape: (32, 224, 224, 3)

Label batch shape: (32, 62)

(32, 1001)

ImageNet predictions



1.1.1 Treinamento do classificador

Agora, carregamos uma versão do Inception V3 sem a camada final de predição, para podermos treinar no novo dataset e obtermos a classificação desejada

```
[10]: feature_extractor_url = "https://tfhub.dev/google/tf2-preview/inception_v3/
      ↪feature_vector/4" #@param {type:"string"}
      feature_extractor_layer = hub.KerasLayer(feature_extractor_url,
                                              input_shape=(224,224,3))
      feature_batch = feature_extractor_layer(image_batch)
      print(feature_batch.shape)
      feature_extractor_layer.trainable = False
      model = tf.keras.Sequential([
          feature_extractor_layer,
          layers.Dense(image_data.num_classes, activation='softmax')
      ])
      model.summary()
```

(32, 2048)

Model: "sequential_1"

Layer (type)	Output Shape	Param #
keras_layer_1 (KerasLayer)	(None, 2048)	21802784
dense (Dense)	(None, 62)	127038

Total params: 21,929,822

Trainable params: 127,038

Non-trainable params: 21,802,784

Agora, nós treinamos o modelo com o dataset Chars74K

```
[11]: predictions = model(image_batch)
      model.compile(
          optimizer=tf.keras.optimizers.Adam(),
          loss='categorical_crossentropy',
          metrics=['acc'])

      class CollectBatchStats(tf.keras.callbacks.Callback):
          def __init__(self):
              self.batch_losses = []
              self.batch_acc = []

          def on_train_batch_end(self, batch, logs=None):
              self.batch_losses.append(logs['loss'])
              self.batch_acc.append(logs['acc'])
              self.model.reset_metrics()

      steps_per_epoch = np.ceil(image_data.samples/image_data.batch_size)

      batch_stats_callback = CollectBatchStats()
```

```
history = model.fit_generator(image_data, epochs=15,
                              steps_per_epoch=steps_per_epoch,
                              validation_data=image_validation,
                              callbacks = [batch_stats_callback])
```

Epoch 1/15

181/182 [=====>.] - ETA: 0s - loss: 2.5876 - acc:

0.6250Epoch 1/15

182/182 [=====] - 49s 269ms/step - loss: 2.5887 - acc:

0.4375 - val_loss: 2.2215 - val_acc: 0.4413

Epoch 2/15

181/182 [=====>.] - ETA: 0s - loss: 1.3990 - acc:

0.7188Epoch 1/15

182/182 [=====] - 41s 226ms/step - loss: 1.3976 - acc:

0.6875 - val_loss: 1.9872 - val_acc: 0.4882

Epoch 3/15

181/182 [=====>.] - ETA: 0s - loss: 0.9981 - acc:

0.6562Epoch 1/15

182/182 [=====] - 41s 226ms/step - loss: 0.9974 - acc:

0.7812 - val_loss: 1.8751 - val_acc: 0.5134

Epoch 4/15

181/182 [=====>.] - ETA: 0s - loss: 0.7630 - acc:

0.6875Epoch 1/15

182/182 [=====] - 41s 225ms/step - loss: 0.7631 - acc:

0.7812 - val_loss: 1.8670 - val_acc: 0.5108

Epoch 5/15

181/182 [=====>.] - ETA: 0s - loss: 0.6033 - acc:

0.9375Epoch 1/15

182/182 [=====] - 41s 225ms/step - loss: 0.6027 - acc:

0.8438 - val_loss: 1.8338 - val_acc: 0.5182

Epoch 6/15

181/182 [=====>.] - ETA: 0s - loss: 0.4865 - acc:

0.9062Epoch 1/15

182/182 [=====] - 41s 225ms/step - loss: 0.4876 - acc:

0.7812 - val_loss: 1.7936 - val_acc: 0.5340

Epoch 7/15

181/182 [=====>.] - ETA: 0s - loss: 0.4005 - acc:

0.9688Epoch 1/15

182/182 [=====] - 41s 224ms/step - loss: 0.4015 - acc:

0.8125 - val_loss: 1.8447 - val_acc: 0.5313

Epoch 8/15

181/182 [=====>.] - ETA: 0s - loss: 0.3400 - acc:

0.9375Epoch 1/15

182/182 [=====] - 41s 225ms/step - loss: 0.3402 - acc:

0.9375 - val_loss: 1.8602 - val_acc: 0.5361

Epoch 9/15

```

181/182 [=====>.] - ETA: 0s - loss: 0.2801 - acc:
0.9375Epoch 1/15
182/182 [=====] - 41s 224ms/step - loss: 0.2801 - acc:
0.9062 - val_loss: 1.8797 - val_acc: 0.5355
Epoch 10/15
181/182 [=====>.] - ETA: 0s - loss: 0.2393 - acc:
0.9688Epoch 1/15
182/182 [=====] - 41s 226ms/step - loss: 0.2392 - acc:
0.9688 - val_loss: 1.8614 - val_acc: 0.5377
Epoch 11/15
181/182 [=====>.] - ETA: 0s - loss: 0.2051 - acc:
0.9688Epoch 1/15
182/182 [=====] - 41s 225ms/step - loss: 0.2057 - acc:
0.9375 - val_loss: 1.8617 - val_acc: 0.5503
Epoch 12/15
181/182 [=====>.] - ETA: 0s - loss: 0.1761 - acc:
1.0000Epoch 1/15
182/182 [=====] - 41s 225ms/step - loss: 0.1761 - acc:
1.0000 - val_loss: 1.9790 - val_acc: 0.5255
Epoch 13/15
181/182 [=====>.] - ETA: 0s - loss: 0.1495 - acc:
0.9688Epoch 1/15
182/182 [=====] - 41s 225ms/step - loss: 0.1500 - acc:
0.9375 - val_loss: 1.9131 - val_acc: 0.5466
Epoch 14/15
181/182 [=====>.] - ETA: 0s - loss: 0.1331 - acc:
0.9688Epoch 1/15
182/182 [=====] - 41s 226ms/step - loss: 0.1329 - acc:
1.0000 - val_loss: 1.9098 - val_acc: 0.5382
Epoch 15/15
181/182 [=====>.] - ETA: 0s - loss: 0.1181 - acc:
1.0000Epoch 1/15
182/182 [=====] - 41s 226ms/step - loss: 0.1183 - acc:
1.0000 - val_loss: 1.9586 - val_acc: 0.5503

```

Com o modelo treinado, podemos visualizar os gráficos de como as medidas de Loss e Accuracy se modificaram com os passos do treinamento

```



[12]: plt.figure()
plt.ylabel("Loss")
plt.xlabel("Training Steps")
plt.ylim([0,2])
plt.plot(batch_stats_callback.batch_losses)

plt.figure()
plt.ylabel("Accuracy")
plt.xlabel("Training Steps")
plt.ylim([0,1])

```

```
plt.plot(batch_stats_callback.batch_acc)
```

[12]: [<matplotlib.lines.Line2D at 0x7f79610d7668>]



1.1.2 Testando o classificador

Após todos os passos completos, podemos executar a predição sobre as mesmas imagens que foram classificadas pelo ImageNet anteriormente e vermos os resultados:

```
[13]: class_names = sorted(image_data.class_indices.items(), key=lambda pair:pair[1])
class_names = np.array([key.title() for key, value in class_names])
class_names
predicted_batch = model.predict(image_batch)
predicted_id = np.argmax(predicted_batch, axis=-1)
predicted_label_batch = class_names[predicted_id]
label_id = np.argmax(label_batch, axis=-1)

plt.figure(figsize=(10,9))
plt.subplots_adjust(hspace=0.5)
for n in range(30):
    plt.subplot(6,5,n+1)
    plt.imshow(image_batch[n])
    color = "green" if predicted_id[n] == label_id[n] else "red"
    plt.title(predicted_label_batch[n].title(), color=color)
    plt.axis('off')
_ = plt.suptitle("Model predictions (green: correct, red: incorrect)")
```


Model predictions (green: correct, red: incorrect)



Também, podemos executar a predição em imagens do dataset de validação:

```
[14]: for validation_batch, label_batch in image_validation:
    print("Validation batch shape: ", validation_batch.shape)
    print("Label batch shape: ", label_batch.shape)
    break

predicted_batch = model.predict(validation_batch)
predicted_id = np.argmax(predicted_batch, axis=-1)
predicted_label_batch = class_names[predicted_id]
label_id = np.argmax(label_batch, axis=-1)
```

```
plt.figure(figsize=(10,9))
plt.subplots_adjust(hspace=0.5)
for n in range(30):
    plt.subplot(6,5,n+1)
    plt.imshow(validation_batch[n])
    color = "green" if predicted_id[n] == label_id[n] else "red"
    plt.title(predicted_label_batch[n].title(), color=color)
    plt.axis('off')
_ = plt.suptitle("Model predictions - validation set (green: correct, red: ↪
↪incorrect)")
```

Validation batch shape: (32, 224, 224, 3)

Label batch shape: (32, 62)

Model predictions - validation set (green: correct, red: incorrect)



1.1.3 Exportando o modelo

Por ultimo, precisamos exportar o modelo para podermos utilizarmos posteriormente

```
[0]: t = time.time()

export_path = "./saved_models/{}".format(int(t))
tf.keras.experimental.export_saved_model(model, export_path)

export_path
```